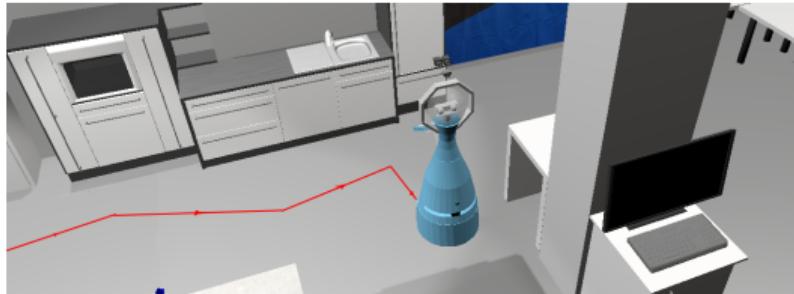


# Learning and Optimizing Probabilistic Models for Planning under Uncertainty in Mobile Robot Navigation



Rob van Bekkum

Delft University of Technology

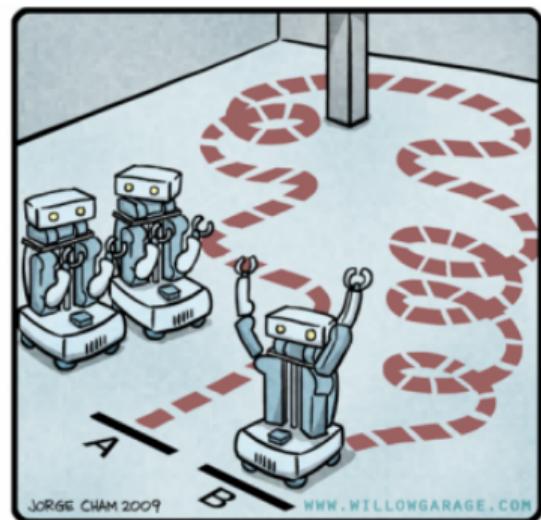
April 5, 2017

# Decision-Theoretic Planning

## DTP Problem Scope

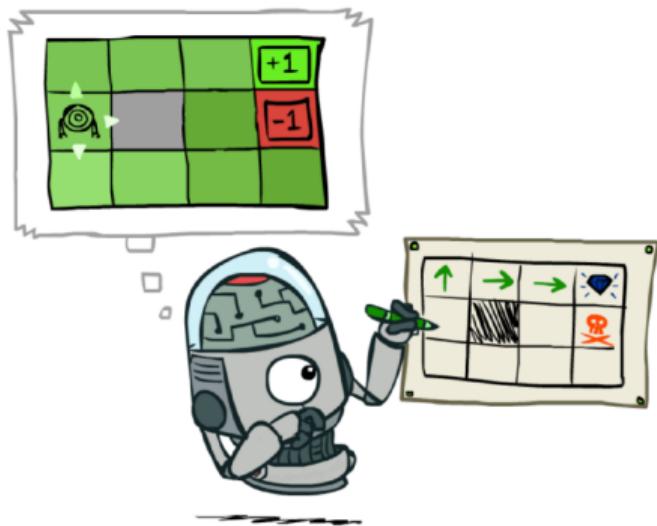
### Systems:

- Dynamics are stochastic, i.e., involve uncertainty
  - ▶ Execution of actions (e.g., robot may slip)
  - ▶ Exogenous events (e.g., doors open/closed)
  - ▶ Observations (e.g., sensor noise)
- Controlled by one or more agents
- Sequential decisions of actions to execute

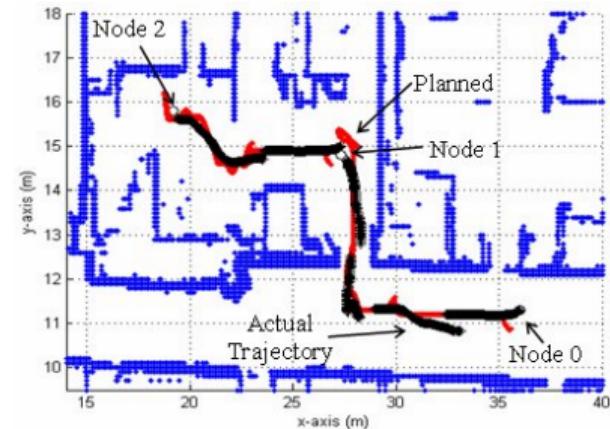


# Decision-Theoretic Planning

## Example Domain: Path Planning



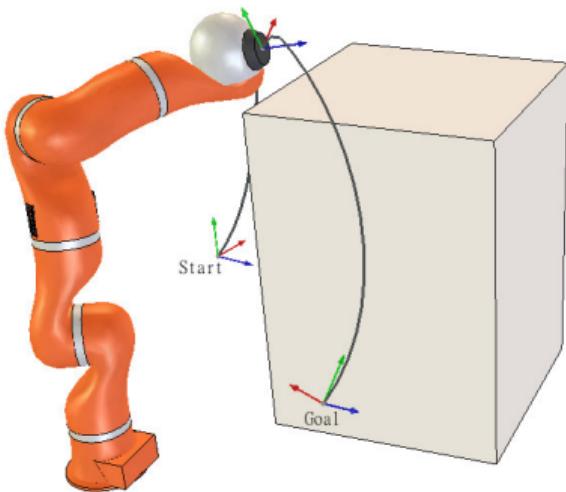
Source: Berkeley CS188 sheets



Source: [http://www.cs.cmu.edu/~gunhee/r\\_psr.html](http://www.cs.cmu.edu/~gunhee/r_psr.html)

# Decision-Theoretic Planning

Example Domain: Motion Planning



Source: <http://www.coppeliarobotics.com/helpFiles/en/motionPlanningModule.htm>

# Decision-Theoretic Planning

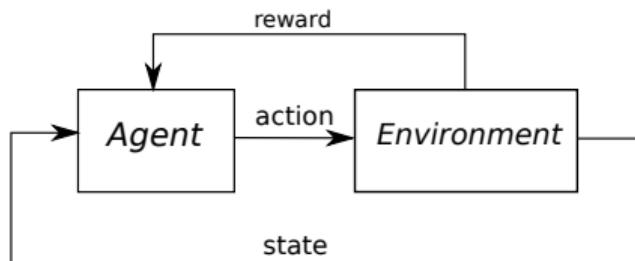
## Markov Decision Processes (MDPs)

In DTP systems are modeled by probabilistic models, e.g. MDPs:

### Definition

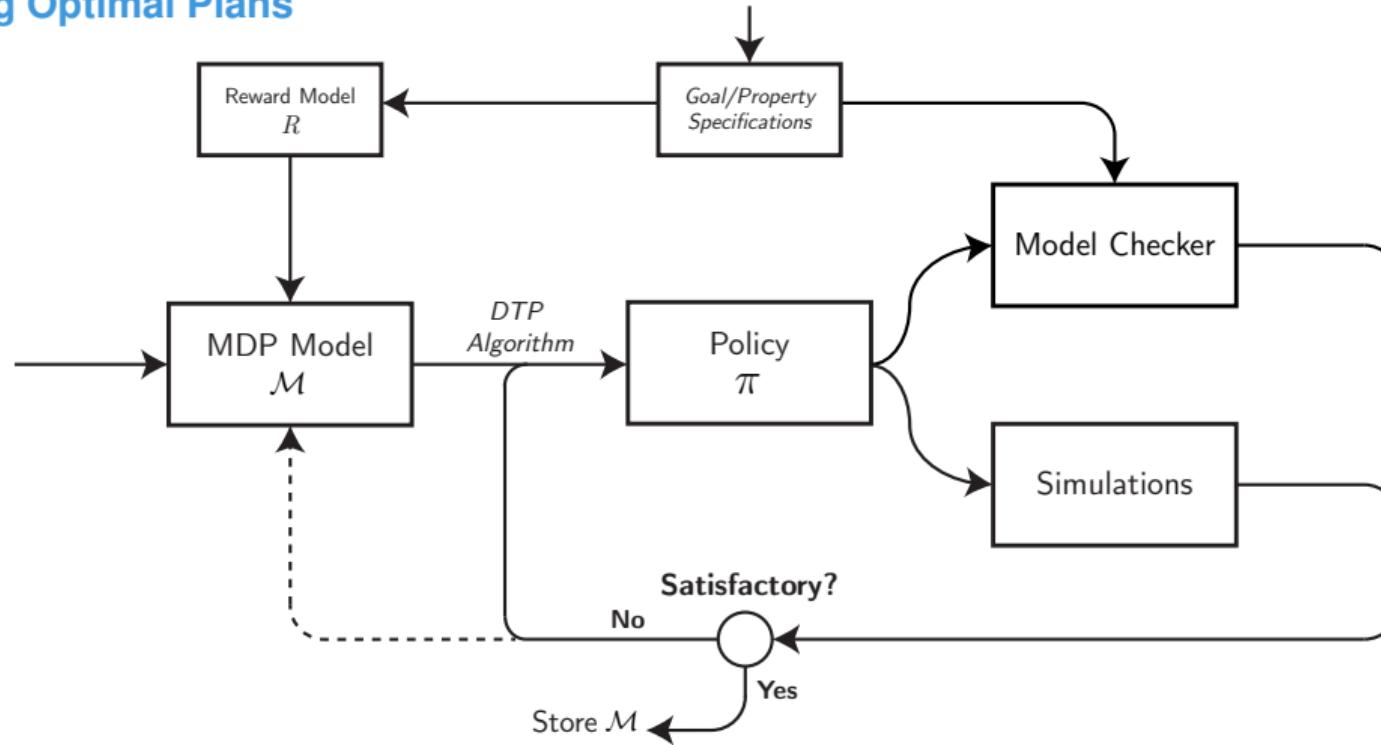
An MDP is a 5-tuple  $\mathcal{M} = (\mathcal{S}, s_0, A, \delta, R)$ :

- $\mathcal{S}$  is the state-space,  $s_0 \in \mathcal{S}$  the initial state
- $A$  is the action-space
- $\delta$  is the transition function (N.B.: accounts for uncertainty)
- $R$  is the reward function (N.B.: defines the goals)



# Decision-Theoretic Planning

## Learning Optimal Plans



# Decision-Theoretic Planning

## Model Development

**Problem:** How to find a suitable MDP model?

**Classical approach:** Model development by a *human designer*, however:

- Requires significant effort (e.g., trial-and-error)
- Typically demands knowledge/experience, accompanied by high costs

**Alternative:** Use Reinforcement Learning instead of Planning, however:

- Requires direct interaction with environment
- One might require *reusable* models, applicable for multiple tasks

# Decision-Theoretic Planning

## Model Development

**Problem:** How to find a suitable MDP model?

**Idea:** Automate the model building process by learning algorithms

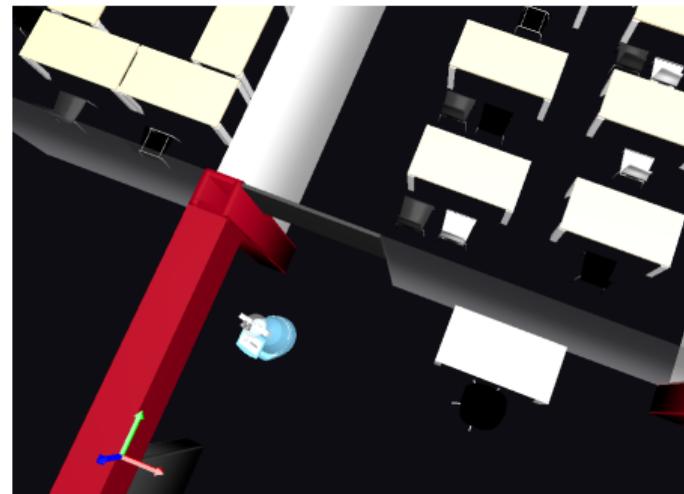
- Learn from (exploration) data about the environment
- Optimization for the best MDP model

# Mobile Robot Navigation

**Goal:** Move robot to area while minimizing traveled distance

MDP Model  $\mathcal{M}$  should define:

- States  $\mapsto$  locations
- Actions  $\mapsto$  robot translations
- Rewards  $\mapsto$  goal areas
- Transition-function



# Model-Learning for Mobile Robot Navigation

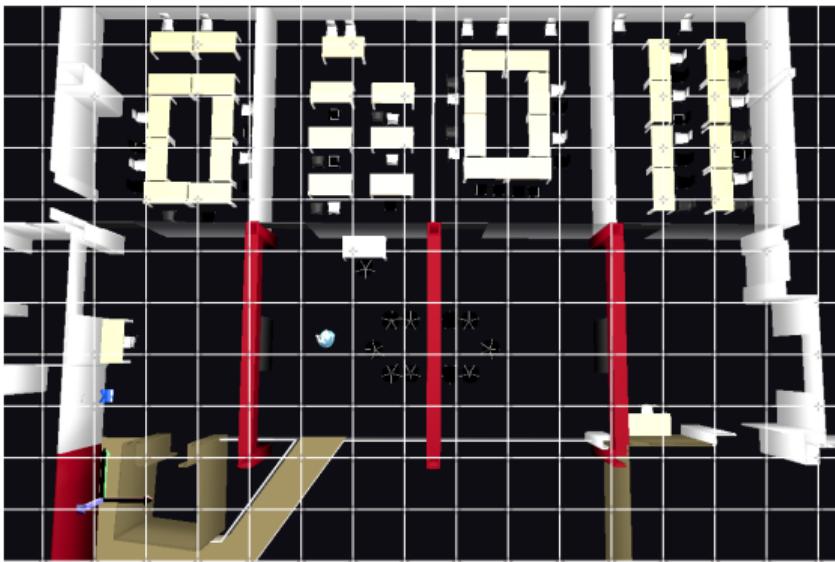
## Learning State-Spaces

**Given:** Execution trace(s) from exploration-stage

### Existing Approaches:

- Grid-discretization
- Best-First Model Merging
- State Merging by Trajectory Clustering

**Issues:** Dimensionality,  
grid resolution



# Model-Learning for Mobile Robot Navigation

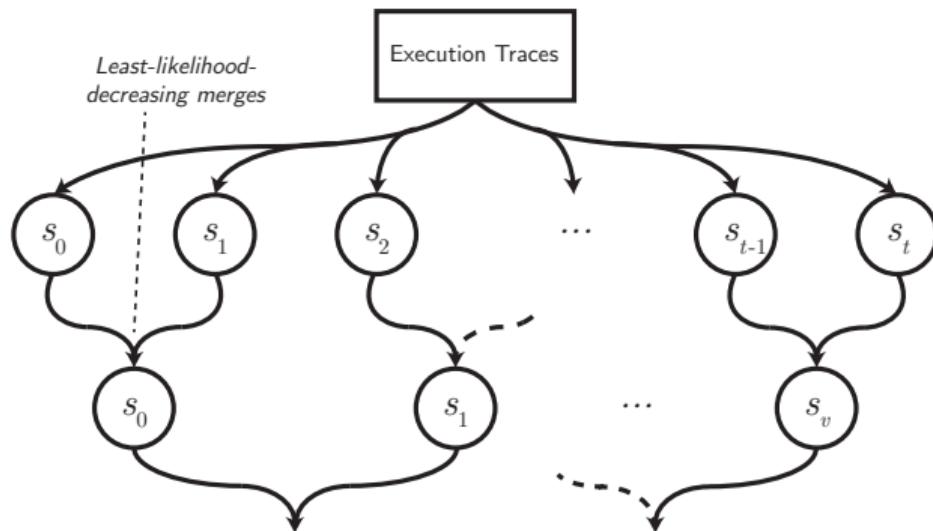
## Learning State-Spaces

**Given:** Execution trace(s) from exploration-stage

### Existing Approaches:

- Grid-discretization
- Best-First Model Merging
- State Merging by Trajectory Clustering

**Issues:** Greedy merging,  
 $O(t^3)$  algorithm



# Model-Learning for Mobile Robot Navigation

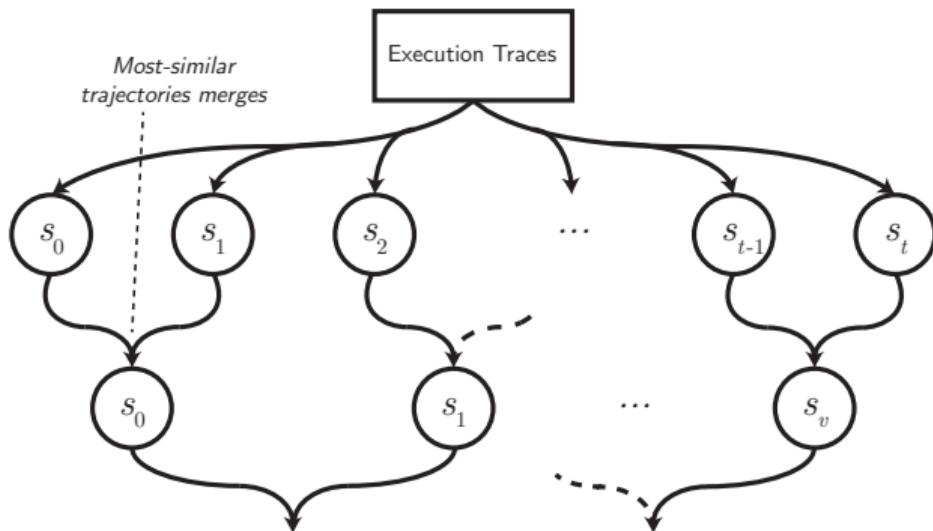
## Learning State-Spaces

**Given:** Execution trace(s) from exploration-stage

### Existing Approaches:

- Grid-discretization
- Best-First Model Merging
- State Merging by Trajectory Clustering

**Note:** For learning POMDPs facing *perceptual aliasing*



# Model-Learning for Mobile Robot Navigation

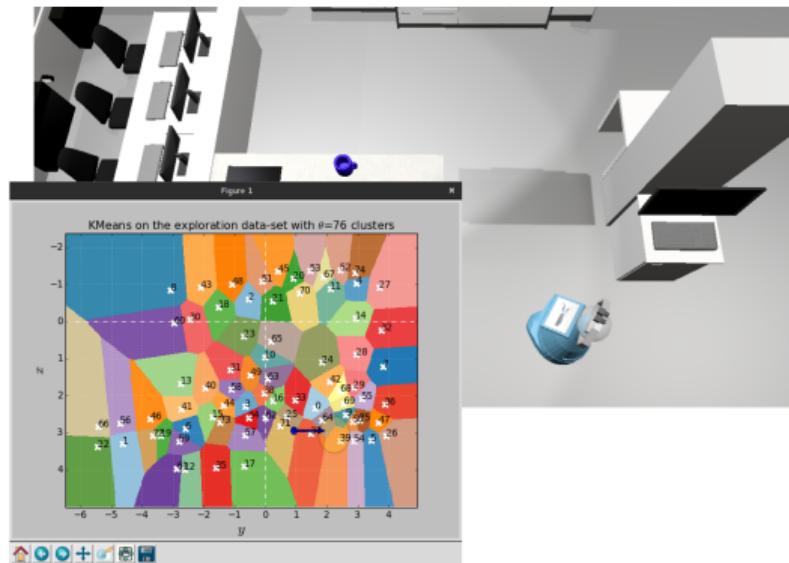
## Learning State-Spaces

**Given:** Execution trace(s) from exploration-stage

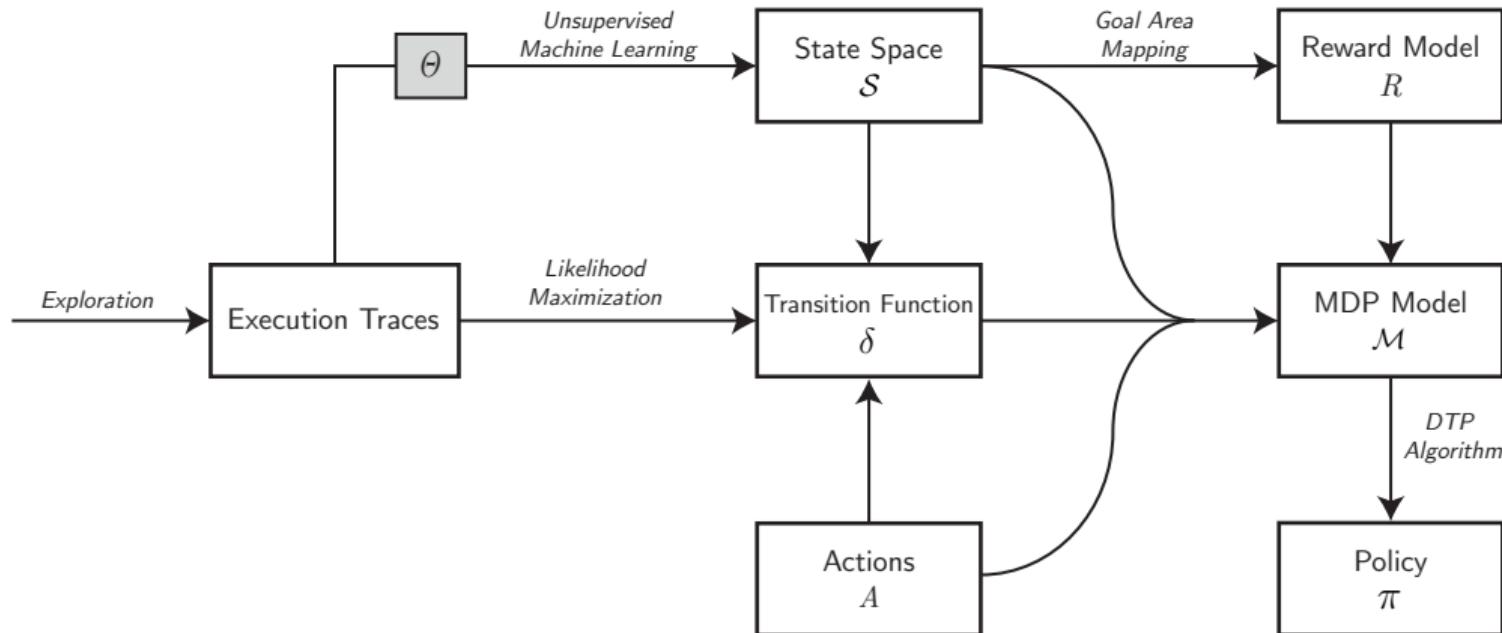
**Approach:** Unsupervised ML  
for state-spaces, i.e.:

- K-Means Clustering
- Gaussian Mixture Models
- ...

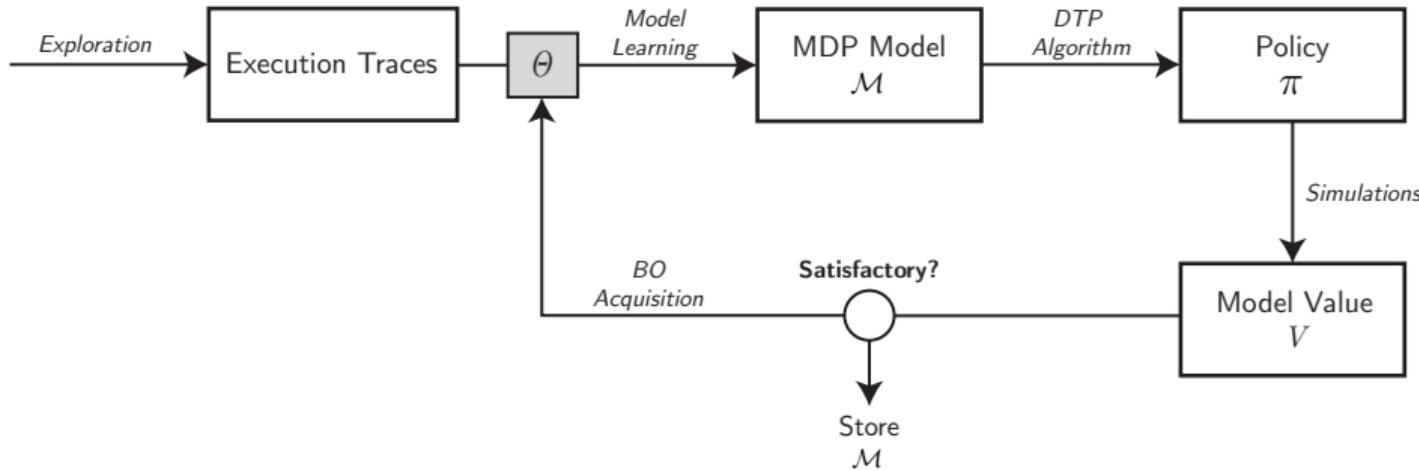
**Variable:**  $\theta$  parameter of  
ML-algorithm



# Model-Optimization Routine



# Model-Optimization Routine



# Bayesian Optimization

- Given an expensive, unknown objective function  $f : \mathcal{X} \mapsto \mathbb{R}$  ( $\mathcal{X} \subset \mathbb{R}^m$ )
- Find global maximizer  $x^* \in \mathcal{X}$  s.t.:

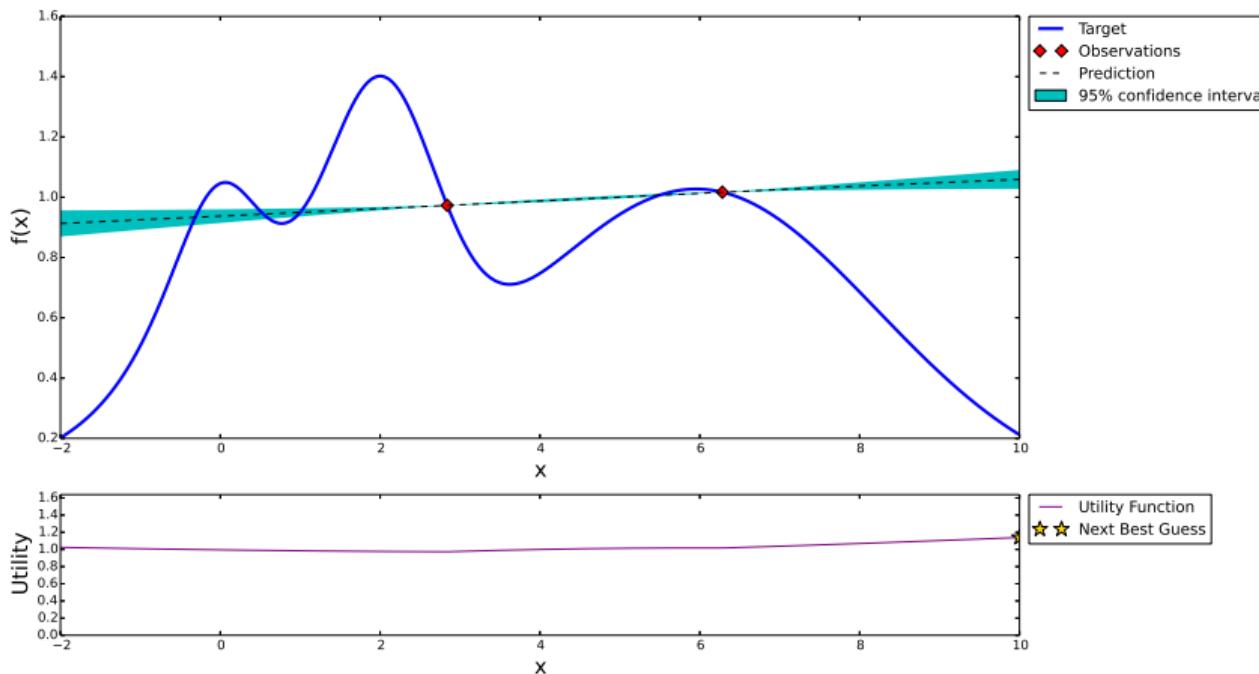
$$x^* = \arg \max_{x \in \mathcal{X}} f(x)$$

- Bayesian Optimization components:
  - ▶ *Prior over  $f$ :* Gaussian Process  $f \sim GP(\mu(\cdot), K(\cdot, \cdot))$
  - ▶ *Evidence set:*  $\mathcal{D}_{1:t} = \{(x_i, y_i) \mid i = 1 \dots t\}$
  - ▶ *Posterior or Surrogate function* ('estimate of  $f$ ')
  - ▶ *Acquisition or Utility function:*  $u : \mathcal{X} \mapsto \mathbb{R}$

# Bayesian Optimization

## Example

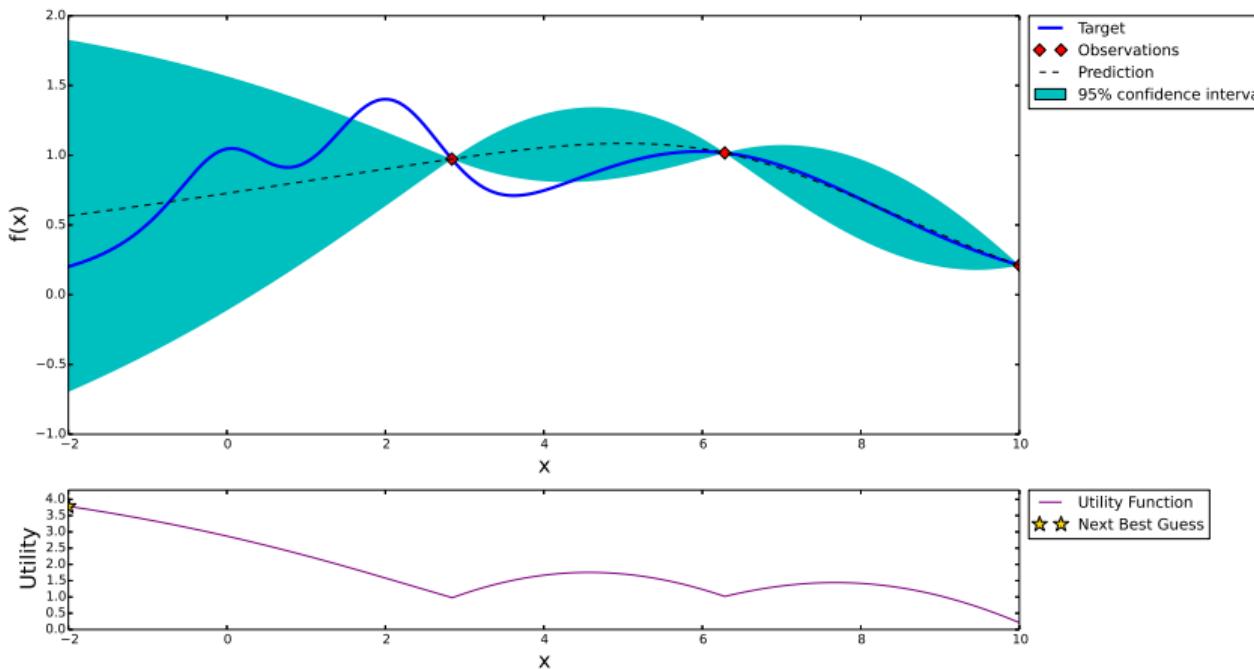
Gaussian Process and Utility Function After 2 Steps



# Bayesian Optimization

## Example

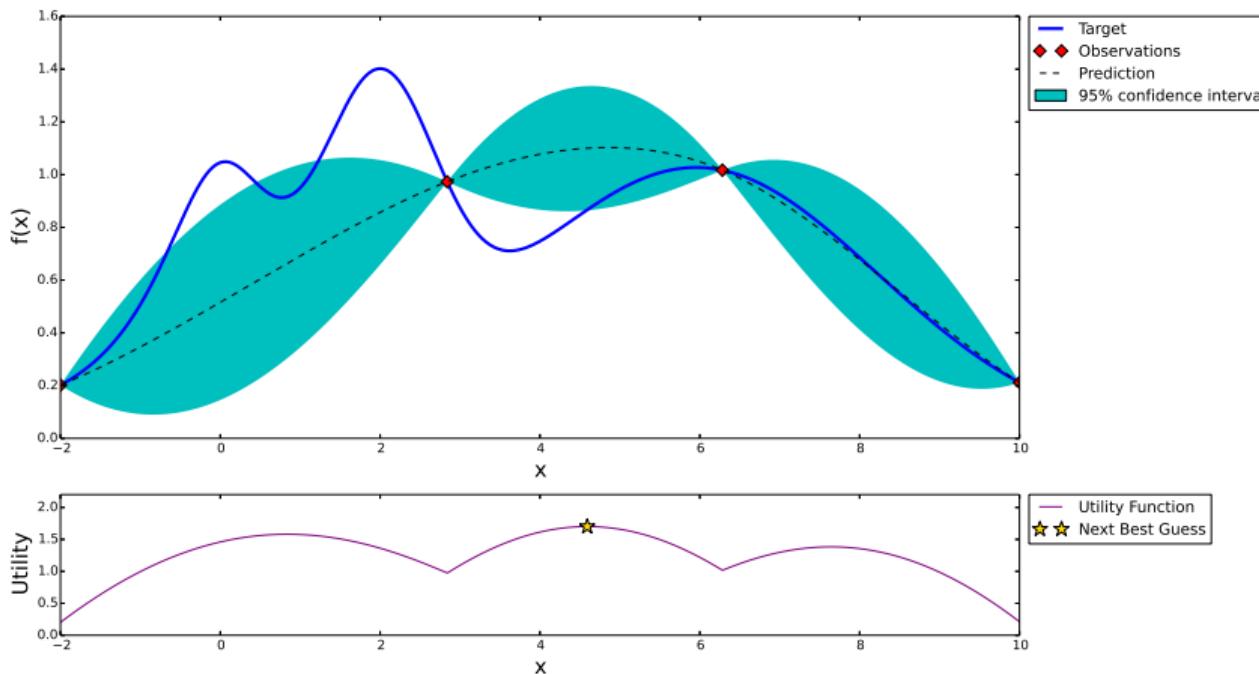
Gaussian Process and Utility Function After 3 Steps



# Bayesian Optimization

## Example

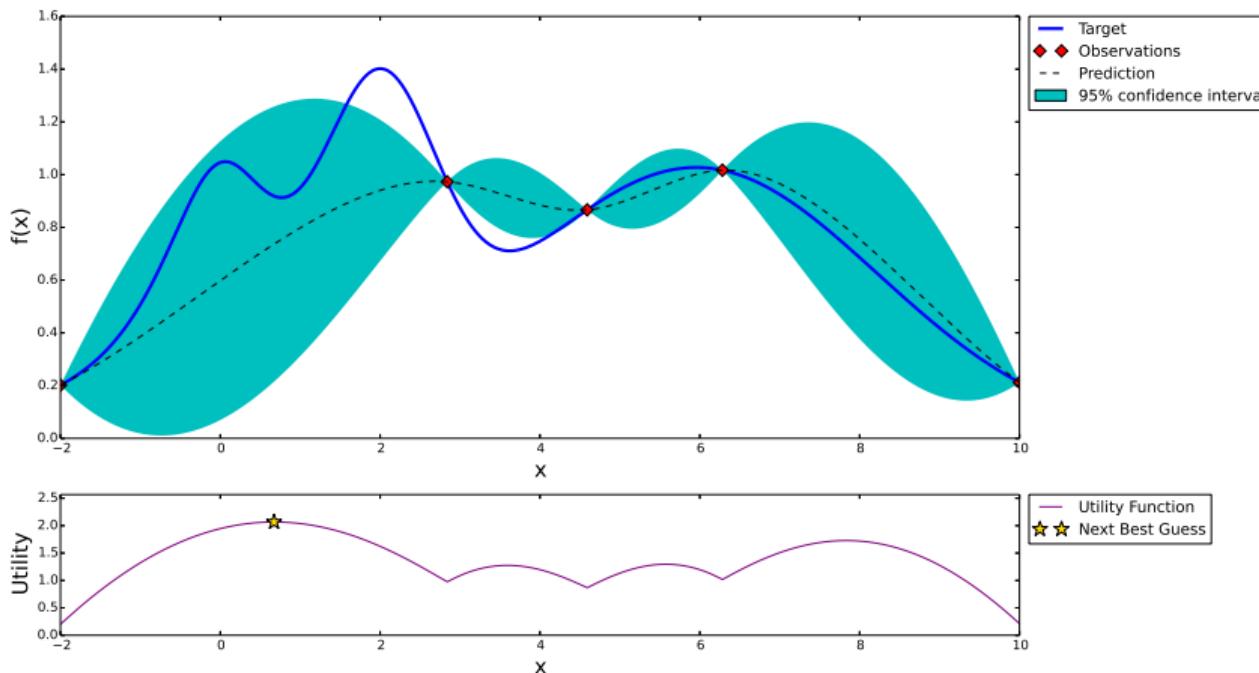
Gaussian Process and Utility Function After 4 Steps



# Bayesian Optimization

## Example

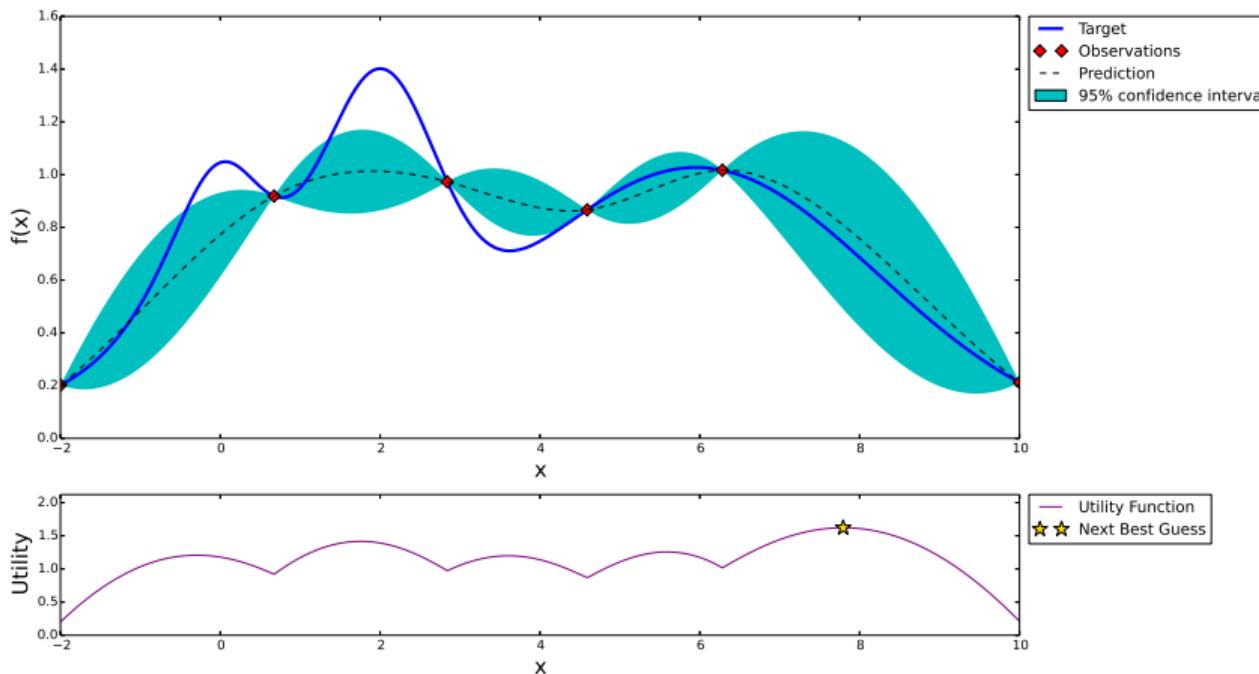
Gaussian Process and Utility Function After 5 Steps



# Bayesian Optimization

## Example

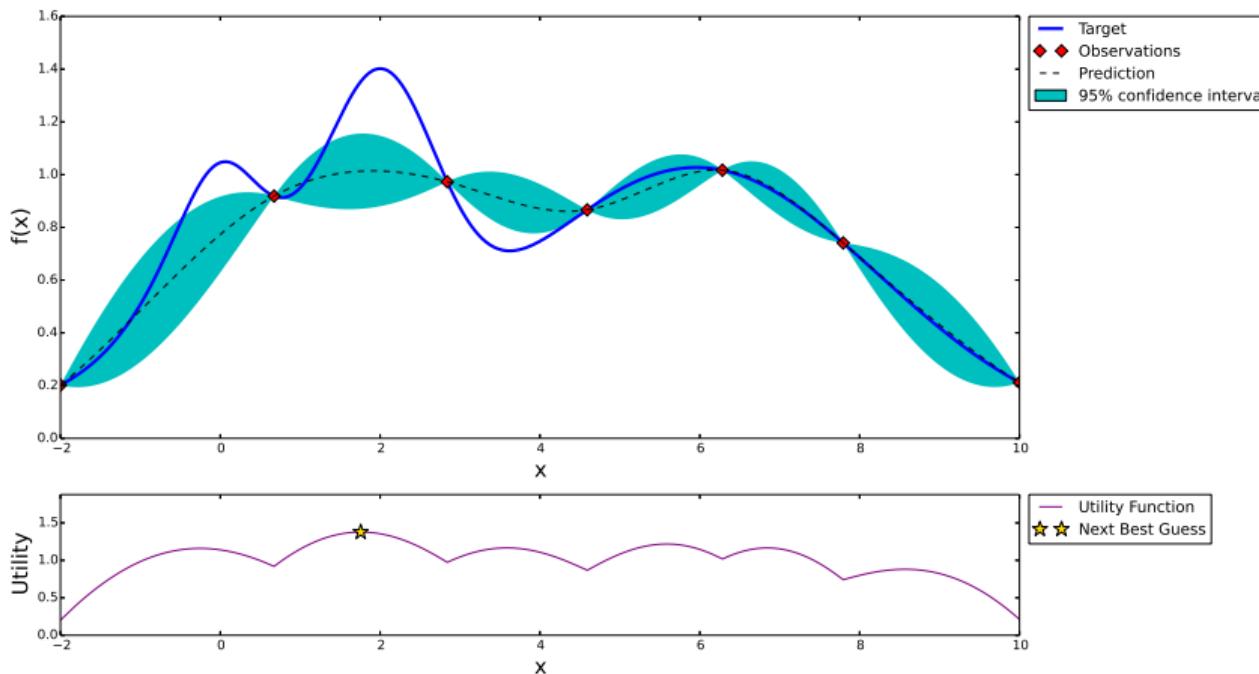
Gaussian Process and Utility Function After 6 Steps



# Bayesian Optimization

## Example

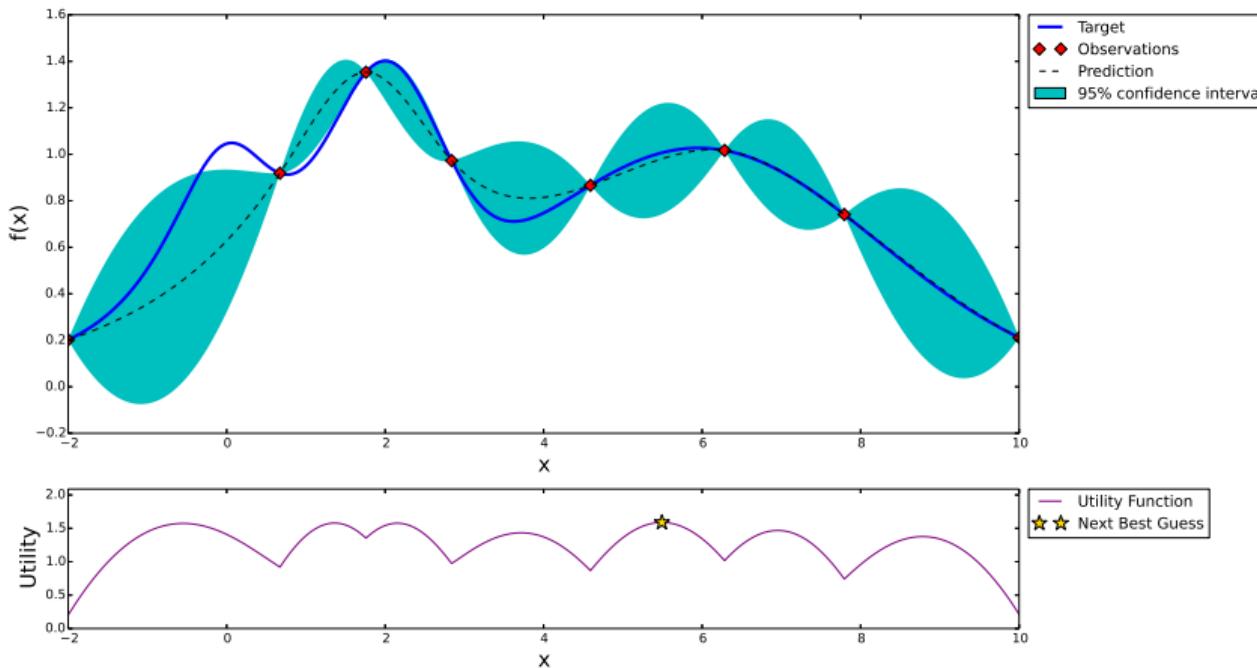
Gaussian Process and Utility Function After 7 Steps



# Bayesian Optimization

## Example

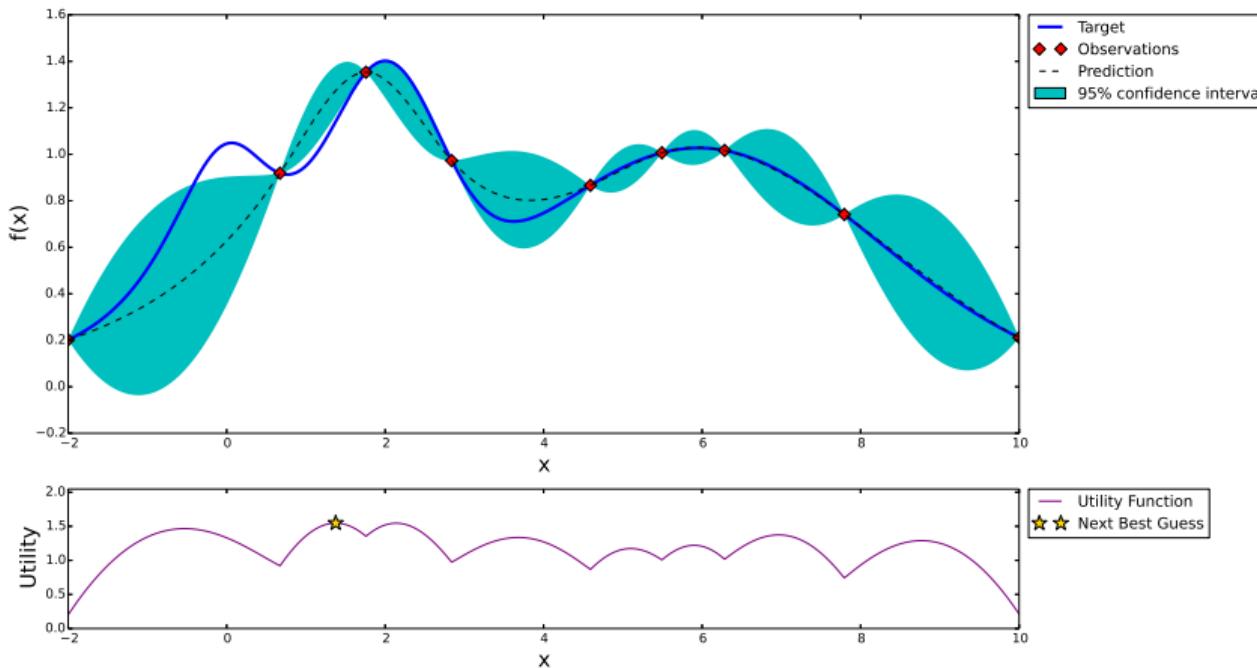
Gaussian Process and Utility Function After 8 Steps



# Bayesian Optimization

## Example

Gaussian Process and Utility Function After 9 Steps



# Implementation of Model-Optimization Routine

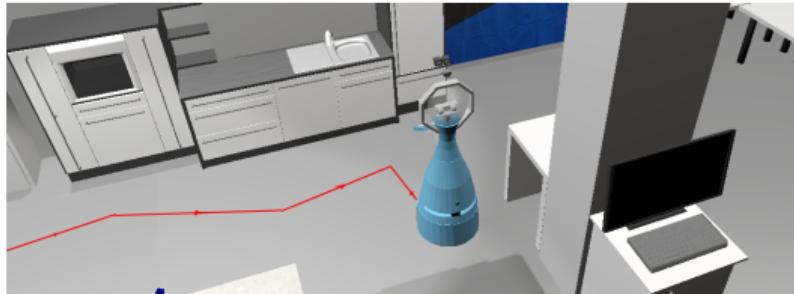
- **Exploration:** 'Randomly' navigating, writing execution trace to files
- **Optimization:** Objective function defined by performance in simulations:
  - ▶ Goal area(s) reached (before time-out)?
  - ▶ Time passed
  - ▶ Part of goal area covered by corresponding final states
- **Software:** Python, ROS, Morse simulator
- **Robot:** Scitos A5, mobile service robot

# DEMO

# Next Steps

- Cost-Sensitive Optimization
  - ▶ BO acquisition based on *Expected Improvement per second*
- Identification of incomplete exploration data
  - ▶ Feedback from simulation trajectories
  - ▶ Use feedback in BO to reduce noise/uncertainty in observations
- Variable Resolution in State-Space Acquisition
- Inspect influence of other variables in the optimization loop
  - ▶ Discount factor  $\gamma$  in MDP solver
  - ▶ ML-algorithm (K-Means, GMM, ...)
  - ▶ Gaussian-Process hyper-parameters?

# Learning and Optimizing Probabilistic Models for Planning under Uncertainty in Mobile Robot Navigation



Rob van Bekkum

Delft University of Technology

April 5, 2017

# References

-  [D. Nikovski and I. R. Nourbakhsh.](#)  
Learning probabilistic models for decision-theoretic navigation of mobile robots.  
In *ICML*, pages 671–678. Citeseer, 2000.
-  [A. Stolcke and S. M. Omohundro.](#)  
Best-first model merging for hidden markov model induction.  
*arXiv preprint cmp-lg/9405017*, 1994.